

Texto pronto pra colar no form de contact-sales da LangChain

URL: <https://www.langchain.com/contact-sales> Preencher os campos abaixo. Forma rápida: copia e cola.

Campo: First name

Fabio

Campo: Last name

Martins

Campo: Work email

diretoria@blinkoftalmologia.com.br

Campo: Message (cola tudo abaixo)

```
Hi LangChain Solutions team,

I run Blink Oftalmologia – an ophthalmology clinic in Brasília, Brazil. We
operate a production WhatsApp scheduling agent ("Lia") built on Anthropic
Claude Sonnet 4.5 + Haiku 4.5, integrating Kommo CRM, Medware (legacy clinic
software), and WhatsApp Cloud API. ~600 active patient conversations/month,
~150 new scheduled appointments/month via the agent.

We've reached a ceiling using raw Anthropic SDK + custom Python FSM. Specific
pain points where I need architectural guidance:

1. Tool-choice enforcement per FSM state – model writes free text instead
  of invoking 'oferecer_slot' when state=AGENDA, even with LIA_TOOLS_ENABLED=1.
2. Reliable production observability – we built 5 custom pillars but a lead
  was abandoned this morning for 90 minutes without any alert firing.
3. Real-transcript evals – we have 700+ pytest with synthetic inputs but
  no eval pipeline against actual production conversations to catch
  regression at scale.

I'd like to evaluate whether migrating to LangGraph + LangSmith Plus (or
Enterprise) is the right next step, or if our current architecture is
sound and just needs polish.

Specific asks:
- 30-45 min call with a Solutions Engineer for architecture sanity-check
- Cost guidance: starting Plus or Enterprise; we're at ~5-10k traces/month
- Pricing for engineering team training (1 dev + 1-2 contractors)
- Whether LangChain Partner Network has consultants experienced with
  healthcare/clinical-scheduling agents in LATAM

I have a 15-page architecture briefing and a 13-page self-audit report ready
to share over NDA. Production health endpoint:
https://blink-agent.6prkfn.easypanel.host/admin/healthz

Best time to talk: weekdays 14:00–18:00 BRT (UTC-3). I prefer English or
Portuguese; happy with either.

Thanks,
Fábio Philipe Martins
Founder / CEO – Blink Oftalmologia
```

Demais campos do form

Campo	Resposta
Job title	Founder / CEO
Company name	Blink Oftalmologia LTDA
Company size	21-100
Company global headquarters	LATAM

Tabela de custos esperados (Plus vs Enterprise)

Item	Developer (grátis)	Plus (\$39/seat/mês)	Enterprise (custom)
Traces/mês incluídos	5k	10k	Customizado
Seats	1	Ilimitado	Customizado
LangSmith Engine (auto-diagnostic)		(\$1.50/LCU)	via créditos
Suporte	Comunidade	Email	SLA + Engenheiro dedicado
Treinamento da equipe			
Architectural guidance			
Hospedagem (Cloud/VPC)	Cloud US/EU	Cloud US/EU	Cloud/Hybrid/Self-hosted
SSO custom			
Faturamento	Self-serve mensal	Self-serve mensal	Invoice anual

Estimativa pro seu caso (600 conversas/mês × ~3-5 traces/conversa = 1800-3000 traces/mês):

- **Plus** já é suficiente em volume de traces, custaria ~\$50-80/mês (1 seat) — bom pra começar
- **Enterprise** só faz sentido se você quiser:
- Self-hosting (dados não saem do Brasil) → exigência LGPD/CFM
- SLA de uptime
- Treinamento estruturado da equipe
- Architectural guidance dedicado

Caminho alternativo se Enterprise for caro: LangChain Partner Network

URL: <https://www.langchain.com/langchain-partner-network>

Partners oficiais (vão te recomendar engenheiros que sabem LangGraph + LangSmith):

1. **Fractional CTOs** que fazem implementação LangGraph profissional
2. **Healthcare-specific consultancies** (raros mas existem)
3. **LATAM partners** (poucos mas crescendo)

Você pode pedir indicação no form acima na seção "Message" — adicione: *"If Enterprise isn't a fit, please connect me with a Partner who can do a paid architecture review + LangSmith setup as a one-off engagement."*

Roteiro pra reunião quando agendarem

Caso eles te liguem (deve ser em 1-2 dias úteis após mandar form):

Primeiros 5 min — você fala: "Briefly: ophthalmology clinic in Brazil, WhatsApp agent built on raw Anthropic SDK + custom FSM, 600 conversations/month. Hit ceiling on tool-calling reliability — model writes free text

instead of invoking structured tools. Need architectural guidance whether to migrate to LangGraph or polish what we have."

Próximos 15 min — você pergunta (use as 5 perguntas do briefing): 1. Tool-choice enforcement por nó do LangGraph 2. Hierarchical agents (router → specialist) — LangGraph ou deepagents? 3. LangSmith Engine sobre traces brutos (sem anotação humana prévia) 4. State persistence across rajada de mensagens (Redis lock atual) 5. Sandboxes pra dry-run agendamento Medware antes de commit

Últimos 10 min — eles propõem: - Plano sugerido (Plus vs Enterprise) - Próximo passo concreto (demo, POC, signing) - Quem é o ponto de contato técnico

Sinais de alerta a observar: - Se eles forcarem Enterprise sem você precisar de SLA/self-hosting → pode estar superdimensionado - Se eles não souberem responder sobre healthcare compliance LGPD/HIPAA → procurar outro consultor - Se eles oferecerem só LangGraph sem LangSmith eval → falta a metade da entrega